



中山大學

操作系统作业

姓名：晋军

学号：18364037

学院：智能工程学院

第一题: Multithreading/Uthread: switching between threads

作业要求: 实现用户的线程切换

实现思路: 像内核线程切换一样, 给线程添加上下文信息用于保存状态, 然后在 thread_schedule 实现交换上下文的功能。

实验过程:

1 修改在 user/uthread.c 中修改上下文结构, 使其能够保存上下文内容, 与 kernel 进程切换的结构类似。

```
struct context {
    uint64 ra;
    uint64 sp;

    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;
};

struct thread {
    char    stack[STACK_SIZE]; /* the thread's stack */
    int     state;              /* FREE, RUNNING, RUNNABLE */
    struct context context;     /* 线程上下文 */
};
```

2 将交换函数添加到 thread_schedule 中, 这里与 kernel 中的 scheduler () 类似, 旧线程为 t, 新线程为 next_thread

```
if (current_thread != next_thread) {
    next_thread->state = RUNNING;
    t = current_thread;
    current_thread = next_thread;
    /* YOUR CODE HERE
     * Invoke thread_switch to switch from t to next_thread:
     * thread_switch(??, ??);
     */
    thread_switch((uint64)&t->context, (uint64)&current_thread->context);
} else
    next_thread = 0;
```

3 在 uthread_switch.S 中具体实现 thread_switch, 这里与 kernel 中的 swithc.S 类似

```

thread_switch:
    /* YOUR CODE HERE */
    sd ra, 0(a0)
    sd sp, 8(a0)
    sd s0, 16(a0)
    sd s1, 24(a0)
    sd s2, 32(a0)
    sd s3, 40(a0)
    sd s4, 48(a0)
    sd s5, 56(a0)
    sd s6, 64(a0)
    sd s7, 72(a0)
    sd s8, 80(a0)
    sd s9, 88(a0)
    sd s10, 96(a0)
    sd s11, 104(a0)

    ld ra, 0(a1)
    ld sp, 8(a1)
    ld s0, 16(a1)
    ld s1, 24(a1)
    ld s2, 32(a1)
    ld s3, 40(a1)
    ld s4, 48(a1)
    ld s5, 56(a1)
    ld s6, 64(a1)
    ld s7, 72(a1)
    ld s8, 80(a1)
    ld s9, 88(a1)
    ld s10, 96(a1)
    ld s11, 104(a1)
    ret /* return to ra */

```

4 创建 thread_create 函数。thread_schedule()在运行切换代码后，希望自动开始运行对应函数，所以在初始化线程的时候，将 ra 寄存器的值赋成对应函数的入口地址，这样在切换结束后运行汇编代码中的 ret 就自动跳转到 ra 指向的位置了；另一个需要初始化的是 sp 寄存器，因为每个线程都各自有一个栈，所以要让自己的 sp 寄存器指向自己的栈底，

```

void
thread_create(void (*func)())
{
    struct thread *t;

    for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
        if (t->state == FREE) break;
    }
    t->state = RUNNABLE;
    // YOUR CODE HERE
    t->context.ra = (uint64)func;
    t->context.sp = (uint64)t->stack + STACK_SIZE;
}

```

运行结果

```

$ usertests sbrkmuch
usertests starting
test sbrkmuch: OK
ALL TESTS PASSED
$

```

第二题：Lock/Memory allocator

题目要求：要求给物理内存分配程序重新设计锁，使得等待锁时的阻塞尽量少。

实现思路：xv6 上只有一个内存链表供应多个 cpu 使用，因此想要实现目标，就希望让每个 cpu 拥有自己的内存链表

实验过程：

1 为每个 cpu 分配内存链表

```

struct {
    struct spinlock lock;
    struct run *freelist;
} kmem[NCPU];

```

2 修改 kinit()和 kfree(), 获取和释放对应 cpu 的锁, 在获取 cpu id 的时候关闭中断, 防止线程切换

```

void
kinit()
{
    for (int i = 0; i < NCPU; i++)
        initlock(&kmem[i].lock, "kmem");
    freerange(end, (void*)PHYSTOP);
}

```

```

void
kfree(void *pa){
    ///
    push_off();
    int id = cpuid();
    pop_off();

    acquire(&kmem[id].lock);
    r->next = kmem[id].freelist;
    kmem[id].freelist = r;
    release(&kmem[id].lock);push_off();
    int id = cpuid();
    pop_off();

    acquire(&kmem[id].lock);
    r->next = kmem[id].freelist;
    kmem[id].freelist = r;
    release(&kmem[id].lock);
}

```

3 最后在添加偷取其他 cpu 空闲内存链表的功能

```

void *
kalloc(void)
{
    struct run *r;

    push_off();
    int id = cpuid();
    pop_off();

    acquire(&kmem[id].lock);
    r = kmem[id].freelist;
    if(r)
        kmem[id].freelist = r->next;
    release(&kmem[id].lock);

    // steal memory
    if (!r)
    {
        for (int i = 0; i < NCPU; i++)
        {
            if (i == id) continue;
            acquire(&kmem[i].lock);
            r = kmem[i].freelist;
            if (r) {
                kmem[i].freelist = r->next;
                release(&kmem[i].lock);
                break;
            }
            release(&kmem[i].lock);
        }
    }

    if(r)
        memset((char*)r, 5, PGSIZE); // fill with junk
    return (void*)r;
}

```

运行结果

```

lock: proc: #fetch-and-add 24315 #acquire() 165760
lock: proc: #fetch-and-add 8715 #acquire() 165830
lock: virtio_disk: #fetch-and-add 8508 #acquire() 75
tot= 0
test1 OK
start test2
total free number of pages: 32496 (out of 32768)
.....
test2 OK
$

```

第三题：Lock/Buffer cache

题目要求：缓解磁盘缓冲区锁竞争情况

实现思路：使用 LRU+hash table 的方法

实验过程：

1 使用 13 作为 hash table 的大小，将 buf 分段映射

```

struct {
    struct spinlock lock;
    struct buf buf[NBUF];           // block[30]
    /** | 循环双向链表 */
    // Linked list of all buffers, through prev/next.
    // head.next is most recently used.
    // struct buf head;

    struct buf buckets[NBUKETS];
    struct spinlock bucketslock[NBUKETS];
} bcache;

```

2 修改和 head 有关的 binit, bget, brelse, bpin, bunpin 函数

```
void
binit(void)
{
    struct buf *b;
    /** 在head头插入b */
    initlock(&bcache.lock, "bcache");
    for (int i = 0; i < NBUKETS; i++)
    {
        initlock(&bcache.bucketslock[i], "bcache.bucket");
        bcache.buckets[i].prev = &bcache.buckets[i];
        bcache.buckets[i].next = &bcache.buckets[i];
    }

    for (b = bcache.buf; b < bcache.buf + NBUF; b++)
    {
        int hash = getHb(b);
        b->time_stamp = ticks;
        b->next = bcache.buckets[hash].next;
        b->prev = &bcache.buckets[hash];
        initsleeplock(&b->lock, "buffer");
        bcache.buckets[hash].next->prev = b;
        bcache.buckets[hash].next = b;
    }
}

void
bpin(struct buf *b) {
    int id = hash(b->blockno);
    acquire(&bcache.lock[id]);
    b->refcnt++;
    release(&bcache.lock[id]);
}

void
bunpin(struct buf *b) {
    int id = hash(b->blockno);
    acquire(&bcache.lock[id]);
    b->refcnt--;
    release(&bcache.lock[id]);
}
```

3 bget 函数里, 如果找到相应缓冲区的话, 就返回, 如果没找到, 就去其他 hash 桶里偷个来放自己所属的缓冲区里, 如果别的 hash 桶里没有的话就报错

```

static struct buf*
bget(uint dev, uint blockno)
{
    struct buf *b;

    int id = hash(blockno);

    acquire(&bcache.lock[id]);

    // Is the block already cached?
    for(b = bcache.head[id].next; b != &bcache.head[id]; b = b->next){
        if(b->dev == dev && b->blockno == blockno){
            b->refcnt++;
            release(&bcache.lock[id]);
            acquiresleep(&b->lock);
            return b;
        }
    }

    // Not cached.
    // Recycle the least recently used (LRU) unused buffer.
    for (int j = hash(blockno+1); j != id; j = (j + 1) % NBUCKETS)
    {
        acquire(&bcache.lock[j]);
        for(b = bcache.head[j].prev; b != &bcache.head[j]; b = b->prev){
            if(b->refcnt == 0) {
                b->dev = dev;
                b->blockno = blockno;
                b->valid = 0;
                b->refcnt = 1;

                // 将别的hash桶里缓冲区放入当前id的缓冲区中
                b->next->prev = b->prev;
                b->prev->next = b->next;
                release(&bcache.lock[j]);
                b->next = bcache.head[id].next;
                b->prev = &bcache.head[id];
                bcache.head[id].next->prev = b;
                bcache.head[id].next = b;
                release(&bcache.lock[id]);
                acquiresleep(&b->lock);
                return b;
            }
        }
        release(&bcache.lock[j]);
    }
    panic("bget: no buffers");
}

```


4 实验结果

```
lock: bcache: #fetch-and-add 0 #acquire() 6749
lock: bcache: #fetch-and-add 0 #acquire() 6711
lock: bcache: #fetch-and-add 0 #acquire() 7911
lock: bcache: #fetch-and-add 0 #acquire() 6213
lock: bcache: #fetch-and-add 0 #acquire() 6212
lock: bcache: #fetch-and-add 0 #acquire() 4158
lock: bcache: #fetch-and-add 0 #acquire() 4155
--- top 5 contended locks:
lock: proc: #fetch-and-add 18128630 #acquire() 4217754
lock: proc: #fetch-and-add 2751180 #acquire() 4210304
lock: proc: #fetch-and-add 562367 #acquire() 4210499
lock: proc: #fetch-and-add 353086 #acquire() 4210305
lock: virtio_disk: #fetch-and-add 286571 #acquire() 1319
tot= 0
test0: OK
start test1
test1 OK
$
```

第四题：Lock/Buffer cache

题目要求：添加大文件系统

实验思路：更改索引，即增加一级索引，由原来的一级索引增加为二级索引。

实验流程：

1 将 FSSIZE 更改为 20000

```
#define FSSIZE      20000 // size of file system in blocks
#define MAXPATH     128  // maximum file path name
```

2 修改宏定义以及 dinode 结构体，使其适配二级索引

```
#define NDIRECT 11
#define NINDIRECT (BSIZE / sizeof(uint))
#define MAXFILE (NDIRECT + NINDIRECT+NDOUBLE_INDIRECT)

// On-disk inode structure
struct dinode {
    short type;           // File type
    short major;          // Major device number (T_DEVICE only)
    short minor;          // Minor device number (T_DEVICE only)
    short nlink;          // Number of links to inode in file system
    uint size;            // Size of file (bytes)
    uint addrs[NDIRECT+2]; // Data block addresses
};
```

3 inode 也需要修改

```
// in-memory copy of an inode
struct inode {
    uint dev;             // Device number
    uint inum;            // Inode number
    int ref;              // Reference count
    struct sleeplock lock; // protects everything below here
    int valid;            // inode has been read from disk?

    short type;           // copy of disk inode
    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+2];
};
```


4 修改 bmap 函数, 主要是将第一次索引的值作为第二次索引的目标, 即重复一次索引。

```
static uint
bmap(struct inode *ip, uint bn)
{
    uint addr, *a;
    struct buf *bp;

    if(bn < NDIRECT){
        if((addr = ip->addrs[bn]) == 0)
            ip->addrs[bn] = addr = balloc(ip->dev);
        return addr;
    }
    bn -= NDIRECT;

    if(bn < NDOUBLE_INDIRECT){
        uint bn_level_1 = bn/NINDIRECT;
        uint bn_level_2 = bn%NINDIRECT;

        // Load level-1 indirect block, allocating if necessary.
        if((addr = ip->addrs[NDIRECT+1]) == 0)
            ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;

        //load level-2 indirect block
        if((addr = a[bn_level_1]) == 0){
            a[bn_level_1] = addr = balloc(ip->dev);
            log_write(bp);
        }
        brelse(bp);

        bp2 = bread(ip->dev, addr);
        a2 = (uint*)bp2->data;

        //load data block
        if((addr = a[bn_level_2]) == 0){
            a[bn_level_2] = addr = balloc(ip->dev);
            log_write(bp2);
        }
        brelse(bp2);

        return addr;
    }

    panic("bmap: out of range");
}
```

5 实验结果

```
$ bigfile  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
.....  
wrote 65803 blocks  
bigfile done; ok  
$
```